

# Neural Time-Series Forecasting: Implementation Plan

Factorized Spatiotemporal Encoder with Iterative Refinement

**Target Challenge:** NSF HDR Neural Forecasting Challenge

**Task:** Predict 10 future timesteps of neural electrode recordings given 10 observed timesteps

**Key Constraint:** Cross-session generalization with day-to-day recording drift

**Evaluation Metric:** Mean Squared Error (MSE) on predicted signals

<b>Version</b>	1.0
<b>Target Audience</b>	Solo Developer
<b>Implementation Language</b>	Python / PyTorch
<b>Estimated Development Time</b>	2-3 weeks

# Table of Contents

- 1. Problem Summary and Constraints
- 2. Motivation: Why This Architecture?
- 3. Architecture Overview
- 4. Component Specifications
- 5. Data Pipeline
- 6. Training Strategy
- 7. Implementation Roadmap
- 8. Hyperparameter Reference
- 9. Debugging Checklist
- Appendix: Quick Reference Card

# 1. Problem Summary and Constraints

## 1.1 The Neural Forecasting Task

The NSF HDR Neural Forecasting Challenge presents a multivariate time-series prediction problem using neural electrode recordings from primates performing motor tasks. The core task is to predict future neural activity given a short observation window, with the critical constraint that models must generalize across recording sessions where electrode behavior may drift.

## 1.2 Data Specifications

Dimension	Specification	Notes
Input Shape	$N \times 20 \times C \times 9$	Samples $\times$ Time $\times$ Channels $\times$ Features
Observed Window	Timesteps 1-10	Model input
Prediction Target	Timesteps 11-20	Only feature[0] is scored
Monkey A (affi)	239 electrodes	~1,150 training samples
Monkey B (beignet)	87 electrodes	~860 training samples
Features	9 per electrode	Feature[0] = target; [1:8] = frequency bands

## 1.3 Key Challenges

- **Session Drift:** Test data comes from different recording sessions where electrode responses may have shifted.
- **Limited Data:** With roughly 1,000 samples per monkey, overfitting is a serious risk.
- **Spatial Structure:** The electrodes record from nearby neural populations with correlated activity.
- **Multi-feature Utilization:** The 8 auxiliary frequency-band features provide additional signal.
- **Short Horizon, High Dimensionality:** Predicting 10 steps  $\times$  239 channels = 2,390 values per sample.

## 2. Motivation: Why This Architecture?

### 2.1 Design Philosophy

#### Core Design Principles:

1. **Simplicity with strong inductive biases** over architectural complexity
2. **Robustness to distribution shift** over raw in-distribution performance
3. **Efficient use of auxiliary features** (frequency bands are free signal)
4. **Factorized computation** to keep parameter counts manageable

### 2.2 Why Not Simpler Approaches?

#### *The GRU Baseline Limitations:*

- Ignores spatial structure entirely, treating electrode 1 and electrode 239 as equally related.
- The flattened input dimension ( $239 \times 9 = 2,151$ ) forces large hidden states, increasing overfitting.
- Cannot share learned temporal dynamics across electrodes.

#### *Why Not a Standard Transformer?*

A vanilla Transformer with (timestep, electrode) tokens would have  $10 \times 239 = 2,390$  tokens. Full self-attention is  $O(n^2)$ , meaning ~5.7 million attention computations per sample—expensive and prone to overfitting on limited data.

### 2.3 The Case for Factorized Attention

- **Physical locality:** Neural dynamics are local in time and space.
- **Computational efficiency:**  $O(T^2 \cdot C + C^2 \cdot T)$  vs  $O((T \cdot C)^2)$ —roughly 10x reduction.
- **Implicit regularization:** Limited attention patterns reduce hypothesis space.
- **Interpretability:** Separate temporal and spatial attention weights are easier to debug.

### 2.4 Why Iterative Refinement?

- Start from a simple initialization (e.g., constant continuation).
- Use the encoder to process concatenated observed + predicted sequence.
- Refine predictions based on learned spatiotemporal dynamics.
- Correct early prediction errors that would otherwise propagate.

### 2.5 Avoiding Electrode ID Embeddings

**Critical Design Decision:** This architecture deliberately avoids learned electrode-specific embeddings. While such embeddings would improve training-set performance, they would memorize session-specific characteristics that don't transfer. Instead, electrode identity emerges from the electrode's temporal pattern and correlation structure with other electrodes.

## 3. Architecture Overview

### 3.1 The Big Picture

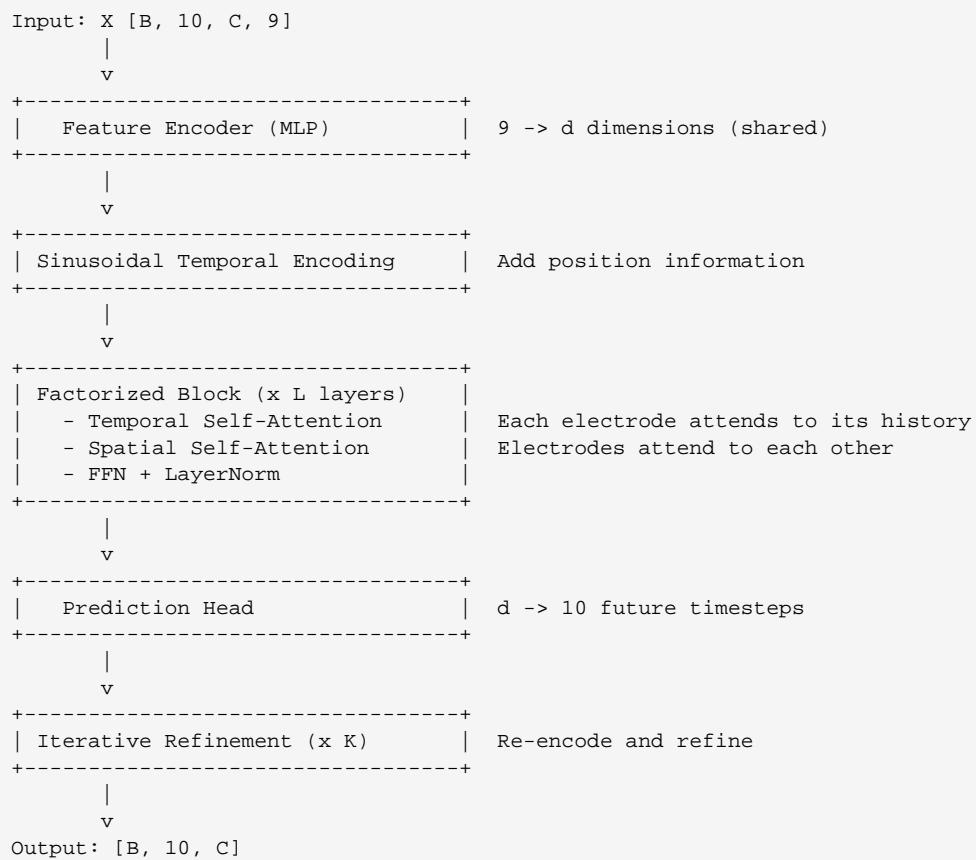
**Stage 1 - Feature Encoding:** Convert each electrode's 9 features into a compact d-dim representation.

**Stage 2 - Spatiotemporal Processing:** Let electrodes "talk to each other" across time and space via factorized attention.

**Stage 3 - Initial Prediction:** Generate a first guess for the next 10 timesteps.

**Stage 4 - Iterative Refinement:** Feed predictions back through the encoder to correct errors ( $K=2-3$  times).

### 3.2 Architecture Diagram



### 3.3 Data Flow Example (Monkey A)

Stage	Shape	Description
Input	[1, 10, 239, 9]	One sample, all features
After Encoder	[1, 10, 239, 64]	Compressed to d=64
Temporal Attn	[239, 10, 64]	Per-electrode sequences
Spatial Attn	[10, 239, 64]	Per-timestep electrode sets
Prediction	[1, 10, 239]	Future timesteps

## 4. Component Specifications

### 4.1 Feature Encoder

The feature encoder transforms 9 raw features per electrode into a d-dimensional representation. It uses a shared MLP across all electrodes and timesteps.

```
class FeatureEncoder(nn.Module):
    def __init__(self, input_dim=9, hidden_dim=32, output_dim=64):
        super().__init__()
        self.encoder = nn.Sequential(
            nn.Linear(input_dim, hidden_dim),
            nn.GELU(),
            nn.LayerNorm(hidden_dim),
            nn.Linear(hidden_dim, output_dim),
            nn.LayerNorm(output_dim),
        )

    def forward(self, x):
        # x: [B, T, C, F] -> [B, T, C, d]
        return self.encoder(x)
```

### 4.2 Positional Encoding

```
class SinusoidalPositionalEncoding(nn.Module):
    def __init__(self, d_model, max_len=50):
        super().__init__()
        pe = torch.zeros(max_len, d_model)
        position = torch.arange(0, max_len).unsqueeze(1).float()
        div_term = torch.exp(
            torch.arange(0, d_model, 2).float() *
            (-math.log(10000.0) / d_model)
        )
        pe[:, 0::2] = torch.sin(position * div_term)
        pe[:, 1::2] = torch.cos(position * div_term)
        self.register_buffer('pe', pe.unsqueeze(0).unsqueeze(2))

    def forward(self, x):
        # x: [B, T, C, d]
        return x + self.pe[:, :x.size(1), :, :]
```

**Why sinusoidal?** Learned positions can memorize dataset-specific patterns. Sinusoidal encodings transfer more reliably to test data with different temporal statistics.

**No spatial positions:** We deliberately omit spatial positional encodings, making spatial attention permutation-equivariant for robustness to electrode drift.

## 4.3 Factorized Spatiotemporal Block

This is the core computational unit. Each block consists of temporal self-attention, spatial self-attention, and a feed-forward network with residual connections.

```
class FactorizedSpatiotemporalBlock(nn.Module):
    def __init__(self, d_model, n_heads, dropout=0.1, ff_mult=4):
        super().__init__()
        # Temporal attention (within each electrode)
        self.temporal_attn = nn.MultiheadAttention(
            d_model, n_heads, dropout=dropout, batch_first=True)
        self.temporal_norm = nn.LayerNorm(d_model)

        # Spatial attention (across electrodes)
        self.spatial_attn = nn.MultiheadAttention(
            d_model, n_heads, dropout=dropout, batch_first=True)
        self.spatial_norm = nn.LayerNorm(d_model)

        # Feed-forward network
        self.ffn = nn.Sequential(
            nn.Linear(d_model, d_model * ff_mult),
            nn.GELU(),
            nn.Dropout(dropout),
            nn.Linear(d_model * ff_mult, d_model),
            nn.Dropout(dropout),
        )
        self.ffn_norm = nn.LayerNorm(d_model)
        self.dropout = nn.Dropout(dropout)

    def forward(self, x, temporal_mask=None):
        B, T, C, d = x.shape

        # === Temporal Attention ===
        x_t = x.permute(0, 2, 1, 3).reshape(B * C, T, d)
        attn_out, _ = self.temporal_attn(x_t, x_t, x_t, attn_mask=temporal_mask)
        x_t = self.temporal_norm(x_t + self.dropout(attn_out))
        x = x_t.reshape(B, C, T, d).permute(0, 2, 1, 3)

        # === Spatial Attention ===
        x_s = x.reshape(B * T, C, d)
        attn_out, _ = self.spatial_attn(x_s, x_s, x_s)
        x_s = self.spatial_norm(x_s + self.dropout(attn_out))
        x = x_s.reshape(B, T, C, d)

        # === Feed-Forward ===
        x = self.ffn_norm(x + self.ffn(x))
        return x
```

## 4.4 Prediction Head

```
class PredictionHead(nn.Module):
    def __init__(self, d_model, n_future=10):
        super().__init__()
        self.head = nn.Sequential(
            nn.Linear(d_model, d_model),
            nn.GELU(),
            nn.Linear(d_model, n_future),
        )

    def forward(self, x):
        # x: [B, T, C, d] - take last timestep
        if x.dim() == 4:
            x = x[:, -1, :, :]
        return self.head(x) # [B, C, 10]
```

## 4.5 Complete Model

```
class NeuralForecaster(nn.Module):
    def __init__(self, n_channels, n_features=9, d_model=64,
                 n_heads=4, n_layers=3, n_future=10,
                 n_refinement_iters=2, dropout=0.15):
        super().__init__()
        self.feature_encoder = FeatureEncoder(n_features, 32, d_model)
        self.pos_encoding = SinusoidalPositionalEncoding(d_model)
        self.blocks = nn.ModuleList([
            FactorizedSpatiotemporalBlock(d_model, n_heads, dropout)
            for _ in range(n_layers)
        ])
        self.pred_head = PredictionHead(d_model, n_future)
        self.n_refinement_iters = n_refinement_iters
        self.n_future = n_future

    def encode(self, x):
        z = self.feature_encoder(x)
        z = self.pos_encoding(z)
        for block in self.blocks:
            z = block(z)
        return z

    def forward(self, x, use_refinement=True):
        B, T, C, F = x.shape
        encoded = self.encode(x)
        pred = self.pred_head(encoded).permute(0, 2, 1) # [B, 10, C]

        if use_refinement and self.n_refinement_iters > 0:
            for _ in range(self.n_refinement_iters):
                pred_features = torch.zeros(B, self.n_future, C, F, device=x.device)
                pred_features[:, :, :, 0] = pred
                full_seq = torch.cat([x, pred_features], dim=1)
                full_encoded = self.encode(full_seq)
                future_encoded = full_encoded[:, T:, :, :]
                pred = self.pred_head(future_encoded.mean(dim=1)).permute(0, 2, 1)
        return pred
```

# 5. Data Pipeline

## 5.1 Normalization Strategy

**Critical Decision:** Use per-sample normalization rather than dataset-global statistics. This is essential for handling session drift, where signal magnitude may change between sessions.

```
class PerSampleNormalizer:
    @staticmethod
    def normalize(x, eps=1e-8):
        # x: [B, T, C, F]
        B, T, C, F = x.shape
        x_flat = x.reshape(B, T * C, F)
        mean = x_flat.mean(dim=1, keepdim=True)
        std = x_flat.std(dim=1, keepdim=True) + eps
        x_norm = (x_flat - mean) / std
        return x_norm.reshape(B, T, C, F), mean, std

    @staticmethod
    def denormalize(x_norm, mean, std):
        return x_norm * std + mean
```

## 5.2 Data Augmentation

```
class NeuralDataAugmentation:
    def __init__(self, electrode_dropout_prob=0.1, noise_std=0.02):
        self.electrode_dropout_prob = electrode_dropout_prob
        self.noise_std = noise_std

    def electrode_dropout(self, x):
        # Randomly zero out electrodes to simulate missing channels
        mask = torch.rand(x.size(0), 1, x.size(2), 1) > self.electrode_dropout_prob
        return x * mask.float().to(x.device)

    def gaussian_noise(self, x):
        return x + torch.randn_like(x) * self.noise_std

    def __call__(self, x, training=True):
        if not training:
            return x
        if random.random() < 0.5:
            x = self.electrode_dropout(x)
        if random.random() < 0.5:
            x = self.gaussian_noise(x)
        return x
```

# 6. Training Strategy

## 6.1 Loss Function

```
class ForecastingLoss(nn.Module):
    def __init__(self, reconstruction_weight=0.3, smoothness_weight=0.05):
        super().__init__()
        self.reconstruction_weight = reconstruction_weight
        self.smoothness_weight = smoothness_weight
        self.mse = nn.MSELoss()

    def forward(self, pred, target, model_outputs=None):
        # Primary loss: prediction MSE
        pred_loss = self.mse(pred, target)
        total_loss = pred_loss

        # Auxiliary: temporal smoothness
        if self.smoothness_weight > 0:
            diff = pred[:, 1:, :] - pred[:, :-1, :]
            smoothness_loss = (diff ** 2).mean()
            total_loss = total_loss + self.smoothness_weight * smoothness_loss

    return total_loss
```

## 6.2 Optimizer Configuration

Use AdamW with weight decay for regularization, combined with cosine annealing learning rate schedule with 10% warmup.

```
optimizer = torch.optim.AdamW(
    model.parameters(),
    lr=1e-3,
    weight_decay=0.01,
    betas=(0.9, 0.98)
)

# Cosine schedule with warmup
total_steps = n_epochs * steps_per_epoch
warmup_steps = int(0.1 * total_steps)

def lr_lambda(step):
    if step < warmup_steps:
        return step / warmup_steps
    progress = (step - warmup_steps) / (total_steps - warmup_steps)
    return 0.5 * (1 + math.cos(math.pi * progress))

scheduler = torch.optim.lr_scheduler.LambdaLR(optimizer, lr_lambda)
```

## 6.3 Early Stopping

Implement early stopping based on validation MSE with patience of 15-20 epochs. Save the best model checkpoint based on validation performance.

# 7. Implementation Roadmap

## 7.1 Week 1: Foundation

- **Day 1-2:** Environment setup, download data, run baseline GRU
- **Day 3:** Implement dataset with per-sample normalization
- **Day 4:** Implement FeatureEncoder and positional encoding
- **Day 5:** Implement FactorizedSpatiotemporalBlock
- **Day 6-7:** Assemble basic model, verify training works

## 7.2 Week 2: Core Features

- **Day 1-2:** Implement iterative refinement
- **Day 3:** Implement multi-component loss
- **Day 4:** Implement data augmentation
- **Day 5:** Training infrastructure (logging, checkpointing)
- **Day 6-7:** Hyperparameter sweep

## 7.3 Week 3: Refinement

- **Day 1-2:** Analyze predictions, identify failure modes
- **Day 3:** Architecture tweaks based on analysis
- **Day 4:** Train ensemble (3-5 models)
- **Day 5:** Package for challenge submission
- **Day 6-7:** Final tuning and submission

## 7.4 Priority Matrix

Priority	Component	Impact
P0	Per-sample normalization	High
P0	Factorized attention	High
P1	Iterative refinement	Medium-High
P1	Electrode dropout augmentation	Medium
P2	Multi-component loss	Medium
P3	Ensemble	Medium

## 8. Hyperparameter Reference

Parameter	Range	Recommended	Notes
d_model	32-128	64	Embedding dimension
n_heads	2-8	4	Must divide d_model
n_layers	2-6	3	Factorized blocks
dropout	0.1-0.3	0.15	Higher for small data
n_refinement	0-3	2	Iterations
batch_size	16-64	32	GPU memory limited
learning_rate	1e-4 to 3e-3	1e-3	With cosine schedule
weight_decay	0.01-0.1	0.01	AdamW

### 8.1 Search Strategy

- **Stage 1:** Fix d\_model=64, n\_layers=3. Search lr in [3e-4, 1e-3, 3e-3] and dropout in [0.1, 0.15, 0.2].
- **Stage 2:** With best LR/dropout, try d\_model in [32, 64, 128] and n\_layers in [2, 3, 4].
- **Stage 3:** Try n\_refinement\_iters in [0, 1, 2, 3].
- **Stage 4:** Fine-tune around best configuration.

# 9. Debugging Checklist

## 9.1 Sanity Checks

- **Overfit one sample:** Can model achieve near-zero loss on a single sample?
- **Gradient norms:** Should be between 1e-5 and 10. Exploding or vanishing = bug.
- **Attention weights:** Temporal attention should show structure, not uniform.
- **Prediction range:** Should match target statistics (mean, std).
- **Ablations:** Removing components should hurt performance.

## 9.2 Common Failure Modes

Symptom	Likely Cause	Solution
Loss stays high	LR too low/high	Try LR finder
Loss NaN	Numerical instability	Add eps, reduce LR
Train good, val bad	Overfitting	More dropout
Predictions constant	Dead model	Check gradients

## 9.3 Expected Performance

Model	Relative MSE
Naive (repeat last)	1.0x (baseline)
GRU baseline	0.7-0.8x
This model (basic)	0.5-0.6x
This model (full)	0.4-0.5x
Ensemble	0.35-0.45x

# Appendix: Quick Reference Card

## Model Architecture

FeatureEncoder(9→64) → SinusoidalPE → [FactorizedBlock×3] → PredHead  
→ Refinex2

## Key Hyperparameters

d\_model=64, n\_heads=4, n\_layers=3, dropout=0.15, n\_refinement=2, lr=1e-3

## Data Flow

Input [B,10,C,9] → Encode [B,10,C,64] → Predict [B,10,C] → Refine → Output

## Critical Decisions

- Per-sample normalization (not global)
- No electrode ID embeddings
- Sinusoidal temporal positions
- Factorized attention (not joint)
- Shared feature encoder

## Loss

Total = MSE(pred, target) + 0.05×SmoothLoss

## Training

AdamW(lr=1e-3, wd=0.01) + Cosine(warmup=10%) +  
EarlyStopping(patience=15)

## Augmentation

ElectrodeDropout(p=0.1) + GaussianNoise(std=0.02)

# Notation Reference

Symbol	Meaning	Value
B	Batch size	32
T	Timesteps	10
C	Channels	87/239
F	Features	9
d	Model dim	64
L	Layers	3
K	Refinement iters	2

**Remember:** Start simple, verify each component, add complexity incrementally.  
The staged approach ensures you always have a working model to fall back on.  
Good luck!